ENCE464 Assignment 2: Computer Architecture

Group 13: Jack Duignan, Isaac Cone, Daniel Hawes

Processor Architecture

This section describes the AMD Ryzen 9 6900HX Central Processing Unit (CPU). This is a laptop CPU based on a 6nm process node. It has eight dual-thread cores for 16 processing units total. It uses the AMD64 (x86-64) instruction set architecture. The overall CPU structure is shown in Figure 1.
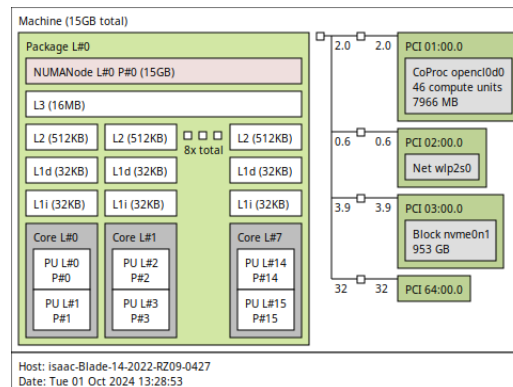

Figure 1: Central Processing Unit (CPU) architecture for the x86-64 AMD Ryzen 9 6900HX.

As mentioned, the cores use the AMD64 architecture, shown in Figure 2 [1]. Each core has several Functional Units (FUs). These include execution units such as Floating Point Units (FPUs), Arithmetic Logic Units (ALUs), memory units and I/O units. By having units specialised for various tasks, the core is able to achieve more than one instruction per clock cycle using pipelining. Each core supports multiple threads as not all functional units can be used simultaneously for example one thread can perform memory access, while the other completes a floating point operation. This is not as performant as two separate cores but does provide an advantage over a single thread.


Figure 2: Typical AMD64 core architecture.

The computer has 15GB of available DRAM which is used for storage of volatile memory. As DRAM slow to access there are multiple levels of higher-speed cache memory implementinga modified Harvard architecture. This cache is split into three layers: the CPU has a shared 16MB L3 cache across all eight cores. Then each core has its own 512 kB L2 Cache and 32kB L1 data and instruction caches. This caching allows both data and instructions to be read in parallel from the shared address space, improving efficiency. Caching is crucial to avoid repeated slow reads from memory.

AMD64 CPUs use the x86-64 instruction set architecture developed by Intel. This is an extension of the ubiquitous x86 architecture that introduces a 64-bit bus while retaining backwards compatibility. The use of a 64-bit architecture allows for much higher addresses of RAM to be read, theoretically up to four exabytes. The x86-64 architecture also has 64-bit general-purpose registers and support for more complicated instructions that allow for more efficient operations on specific data types. An example of this is the Single Instruction Multiple Data (SIMD) instructions that allow operation on vectors of data instead of a single value.

Multithreading

Multithreading was used to improve CPU utilisation. This is achieved by making the program run on multiple cores by separating the Poisson algorithm across multiple worker threads. Each thread is responsible for computing a portion of the cube, divided across slices in the k dimension. The start of a slice is calculated through a simple formula seen below:

$$k_{\text{start}} = 1 + \frac{i\lceil(N-2)\rceil}{t} \tag{1}$$

Where i is the thread number N is the number of nodes and t is the number of threads. The end of a slice is calculated in a similar way as:

$$k_{\text{end}} = (i+1)\frac{\lceil(N-2)\rceil}{t} + 1, k_{\text{end}} = \begin{cases} N-1 \text{ if } k_{\text{end}} > N-1 \\ k_{\text{end}} \quad \text{otherwise} \end{cases} \tag{2}$$

Each worker thread is passed the program variables (`curr`, `next` etc.) and the slice it is assigned. The Von Neumann and inner iterations are then applied as required. To prevent race conditions caused by parallel execution a barrier is used. This is a `pthread_barrier_t` type with a limit equal to the number of worker threads. When all threads have completed an iteration and called the wait function, the barrier allows the program to proceed in sync. After each thread completes its calculations for one iteration, the buffers for the next and current iterations need to be changed. This is done by swapping the pointer addresses, avoiding the need for expensive memory copying operations.

The results of the multithreading implementation can be seen in Figure 3. The results show that the execution time decreases as the number of threads is increased. This expected because the program can make use of multiple cores in parallel. The solution reaches an execution time asymptote at approximately 20 threads. After this point the execution time is near constant. The minimum execution time is reached by 12 threads. This is because the results are captured on an Intel i5 12400f processor which has 6 cores and 12 threads. The most common operation completed by the Poisson implementation is floating point arithmetic which requires a floating point unit to execute efficiently. Two threads are able to utilise a single FPU as it takes time to load memory back and forth. This means that the theoretical maximum number of threads that can complete floating-point operations simultaneously is 12 which is reflected in the minimum computational time value.


Figure 3: A comparison of the poisson solver software execution times with different number of threads used

Caching

The Poisson algorithm requires a large number of memory operations to load the data from the respective arrays and store the result. These memory operations can add significant overhead. This is because the reading/writing from DRAM takes orders of magnitude longer than the execution time of the CPU. Modern CPUs minimise this impact with small high-speed memory caches. These store recently accessed data, minimising costly DRAM access operations. As discussed in Section 1, the largest cache on the analysed CPU is 16MB. This is not enough to store even one of the arrays needed, (for a 401 node cube the arrays are 515MB). This means the cache can only store small sections of memory. Optimising cache utilisation can signiciantly speed up program execution.

The largest number of memory operations in the code occurs during the computation of the inner nodes of the cube. This uses three nested `for` loops that iterate over the layers, columns, and rows. These iterations can be executed in any order to achieve the correct result. One way to ensure optimal cache utilisation is to consider spatial locality when accessing memory. By accessing neighbouring data consecutively, required data will almost always be located in the cache for each node. To achieve this, the optimal iteration scheme is for the inner loop to iterate along the i-axis (a row), as it increments by a single item in memory each iteration. This is based on data shown in Figure 4. As seen in the plot, iteration schemes where the inner most loop is the i-axis outperform any other scheme by a significant margin. The execution time is directly proportional to what iteration level the i-axis is performed. The i, k, j iteration scheme is five times slower than the optimal k, j, i scheme.


Figure 4: A comparison of execution time for various cube sizes with a range of iteration schemes.

The performance differences in iteration schemes can be explained by examining the number of cache misses occurring during execution. If the CPU needs to access data that is not in the cache, this is a miss, and the CPU must then read this data from DRAM. Table 1 compares the L1 cache read and write miss rates for the same program using the best (k, j, i) and worst (i, k, j) performing iteration schemes respectively. This data was generated using the `cachegrind` program. The total number of memory accesses is equal, but the optimal iteration method has approximately five times fewer misses, indicating significantly improved cache utilisation hence the lower execution time.

Table 1: Memory access cost at various cache levels for the AMD Ryzen 6900HX.

Iteration	Dmr	Dmr %	Dmw	Dmw %
i, k, j	908,929,991	5.03	296,704,937	96.7
k, j, i	168,924,763	0.93	42,754,806	14.0

Profiling

Profiling was used throughout all stages of this project's development to identify which areas of the program would benefit from optimisation. Amdahls' law was applied to select appropriate areas for optimisation. This meant that areas of code which are most often used were optimised first, as this would have the largest impact on performance. To make profiling easier the various components of the code were compartmentalised into functions. While this does add some execution time (due to stack overheads) it allows the profiling tool `gprof` to provide more granular results.

Profiling was conducted on both optimised and non-optimised code to gain a holistic understanding of the program's execution. A breakdown of the execution times and call counts for a non-optimised run of the program with a 201-node cube over 300 iterations using 20 threads can be seen in Table 2. This is compared to using debug optimisation (-Og) on the same cube size in Table 3.

Table 2: GProf results for a non-optimised run of the program with 201 nodes 300 iterations and 30 threads.

Function	Percentage	Call Count	Time per call (ms)
`poisson_iteration_inner_slice`	96.47%	2251	24.05
`apply_von_neuman_boundary_slice`	3.53%	2222	0.91
`wait_to_copy`	0%	2400	0
Setup	0%	1	0

Table 3: GProf results for an Og optimised run of the program with 201 nodes 300 iterations and 30 threads.

Function	Percentage	Call Count	Call Time (ms)
`poisson_iteration_inner_slice`	93.45%	2269	11.45
`apply_von_neuman_boundary_slice`	6.55%	2275	0.80
`wait_to_copy`	0%	2389	0
Setup	0%	1	0

The results found in Table 2 and 3 show that in both runs the largest time cost is the iteration over the inner slice of the cube. This is expected as it performs the majority of the floating point operations. As expected the compiler optimisations have reduced the iteration time by more than half. Interestingly, the Von Neumann boundary condition execution time was only reduced by 12%. This may be due to the significant number of conditional checks required by this function which cannot be optimised.

In earlier versions of the program, the Von Neumann boundary was called at every inner loop of the main Poisson iteration. Based on profiling it was identified that this was as a bottleneck as it is unnecessary to call this for all the inner nodes. This was moved to its own self-contained iteration that only iterates over the outside nodes. This significantly reduces the number of conditional checks needed thus reducing the execution time as there are fewer instructions per iteration.

Another example of profiling helping in the optimisation of code is the barrier wait that is used to synchronise the threads. Originally it was hypothesised that these waits would greatly increase the execution time as threads take different amounts of time to complete due to CPU allocation and allocation of different areas of the cube. By profiling the code with these barriers implemented it was discovered that the barrier wait does not add appreciable execution time. This is shown in Table 2.

Compiler Optimisation

Modern compilers, particularly for the C Programming Language, are extremely well optimised. This makes it near-impossible for a programmer to "beat" the performance of compiler-generates assembly code. Compiler optimisation modifies the standard operation of a compiler to produce assembly code optimised for some specific purpose, usually execution time or program size. These optimisations result in tradeoffs. For example, a program optimised for execution speed may be larger than an unoptimised program. Because of this, compiler optimisation are disabled by default, and are enabled using compiler flags. The compiler flags are designed to enable a range of specific optimisations, and are organised to focus on different optimisation strategies. The flags, their strategies, and the resulting compiled size for the Poisson implementation are sumarrised in Table 4.

Table 4: The impact of optimisation flags on the program size of the Poisson algorithm implementation.

Flag	Optimisation Strategy	Size (kB)
-O0	no optimisation performed	68.0
-O1	performance/size without adding compilation time	54.5
-O2	performance/size, can increase compilation time	54.3
-O3	performance, can increase size and compilation time	51.4
-Ofast	O3 with impact to floating point math precision	67.9
-Os	optimise for size over all other factors	45.8

Each of the optimisation strategies was applied to the Poisson algorithm implementation. Using program size to predict performance impacts, it is expected that all of the performance-focused optimisation methods will result in a shorter execution. The degree of improvement should theoretically increase as the methods become more focussed on performance with less regard for compile time or program size. The results for each method are compared in Figure 5.


Figure 5: Comparing a range of performance optimisation options for the Poisson algorithm.

The results show that any performance optimising method makes a significant difference to execution time, particularly at larger cube sizes. What was not expected is that all optimisation methods perform approximately equally. This means that the additional space-intensive optimisations enabled by -O3 are not able to be applied to the Poisson algorithm. One reason for this is that the size of the main loops of the program are defined at runtime, preventing compiler loop unrolling.

Individual Topic 1: Branch Prediction (Jack Duignan)

Execution on a CPU is optimised using pipelining which allows multiple instructions to be executed in parallel on a single core. It does this by loading the next set of instructions while the previous set are being executed. When branch instructions occur a control hazard develops as the CPU cannot predict which set of instructions to load. This can cause costly pipeline flushes if the CPU loads the wrong instructions. To mitigate the number of times this flush occurs the CPU predicts which path execution will follow before the conditional branch is executed. This is branch prediction.

In the Poisson algorithm software, there are several control hazards caused by conditional branching. One of the most significant of these is during the inner iteration of the Poisson equation. In this operation, the code uses several conditional instructions to check the location of the current node and apply the correct iteration. This results in the implementation of the iteration having 12 conditional jump instructions per node update (if complied without optimisations). To improve the speed of the implementation it was decided to reduce the number of these jumps where possible to reduce the amount of branch prediction required by the CPU.

The reason the implementation required so many conditional instructions is that it uses the same nested `for` loop to iterate over both the Von Neumann boundary (the Dirichlet boundary can be applied once during initialisation) and the inner nodes of the cube. It was hypothesised that moving these out of the same loop and reducing the Von Neumann iteration to only over the outer nodes would reduce the number of condition branch control hazards per iteration. This was achieved by splitting each iteration into two components first the Von Neumann boundary is applied to only the nodes required then the nested loop only updated the inner nodes. This change completely removed the conditional instructions in the main iteration loop (which is called most often) removing the largest control hazard from the program.

A comparison of the software with and without conditional control hazards can be seen in Figure 6. This shows that the program's execution has been reduced by 10%. With the real benefits occurring at larger cube sizes as more iterations mean more conditional branch issues. This result is expected as by reducing the number of conditional branches the CPU can optimise the use of pipelining as the number of possible flushes is reduced. This change also has the added benefit of reducing the number of instructions per iteration. This is due to the Von Neumann boundary application only iterating over nodes that it will need to be applied to as opposed to all nodes in the cube. Crucially this time gap still appears when optimisation is applied showing that the problem cannot be solved by the compiler.


Figure 6: A comparison of the poisson solver execution times for different branch schemes.

Individual Topic 2: GPU Implementation (Isaac Cone)

Graphics Processing Units (GPUs) are specialised processors with many cores optimised for performing repeated operations. GPUs are significantly faster than CPUs in some applications due to massively parallel execution, a much higher memory bandwidth, and greater instruction throughput. This section will discuss how the NVIDIA 3070 Ti laptop GPU can be leveraged to enhance the performance of the Poisson algorithm. The 3070 Ti architecture shown in Figure 7 consists of 48 Streaming Multiprocessors (SM) each with 128 cores for a total of 6144. These cores run up to eight threads each. The hardware executes a custom kernel function using the Compute Unified Device Architecture (CUDA) API [2]. The API uses variably sized blocks of threads and automatically handles the allocation of threads and blocks to the hardware. The poisson algorithm involves a single repeated operation, making it ideal for implementation on the GPU. It is expected that the 3070 Ti, which has 6144 CUDA cores will significantly outperform a CPU, particularly on larger cube sizes.


Figure 7: Multiprocessor and CUDA core architecture of the NVIDIA 3070 Ti GPU.

To simplify implementation, the Poisson algorithm was reduced to a single kernel function, no longer optimised. To leverage the GPU VRAM bandwidth, computations are performed entirely in the GPU VRAM. This has the limitation that the data size cannot exceed the GPU VRAM capacity, 8GB in this case. For execution on the 3070 Ti, the threads per block were set to 512 in an 8-thread cube. The gridSize cube of blocks is then dynamically sized for a given N by Equation 3.

$$\text{gridSize} = (N + \text{blockSize} - 1)/\text{blockSize} \tag{3}$$

Figure 8 compares the CUDA program to the optimal CPU program. For a 101 cube, the CPU outperforms the GPU. This is because the overhead from copying to VRAM exceeds the benefit of increased thread count. As cube sizes grow, the GPU significantly outperforms the CPU. This is because the advantage of massively parallel execution becomes more pronounced with more computations. Above a 601-sized cube, memory exceeds 8GB, preventing execution on the 3070 Ti. This could be avoided by batching data into smaller sections, but this would add overhead. Additionally, the CUDA program was written as a proof of concept, meaning there are optimisations that could realise significant performance improvements. One bottleneck is the many conditional statements in the kernel. The GPU uses a shared program counter for groups of threads to minimise overhead. This means that any differences in thread execution forces faster threads to wait for slower ones.


Figure 8: CUDA implementation of the poisson algorithm compared to optimal CPU solution.

Individual Topic 3: SIMD Implementation (Daniel Hawes)

Single Instruction, Multiple Data (SIMD) is a technique used to perform the same operation on multiple data points simulataneously. This is particularly useful in applications where the same operation is performed over large data sets, such as large 3D arrays. This section will discuss how SIMD can optimise the poisson algorithm. SIMD uses the Advanced Vector Extension (AVX) instruction set. The results of the implementation are from tests on the AMD Ryzen 7 4700U CPU.

The AVX2 instruction set allows for 256-bit wide registers to be used, which can hold four 64-bit double precision floating point numbers. This can be applied to the poisson algorithm by performing the same operation on 4 nodes simultaneously. This is done across four nodes in the i dimension. It does one calculation in steps of 4 with a calculation manually performed at the end if the inner slice i dimension nodes are not divisble by 4. This finishes off the remaining nodes.

Between `i` and `i+3` load each node surrounding the four calculated nodes can be loaded into a vector register. This results in 6 different vector registers holding values for the calculation. Since all the vectors will use the add operation, the add operation can be done simultaneously for the 4 nodes and their respective surrounding nodes. It can be hypothesised that there should be approximately a 4x increase on the inner slices of the cube since it can now calculate 4 nodes at once. With this hypothesis in mind, the results of the SIMD implementation can be seen in Figure 9. These results come from SIMD only implemented in the inner slice calculation since SIMD is only very useful when the same operation is performed on multiple data points multiple times.


Figure 9: Comparison with SIMD implemented in the inner slice and without SIMD.

The results show that despite the implementation of SIMD in the inner slice layers, the execution time is decreased. Based on SIMD principles, this conclusion didn't seem correct, so after profiling the SIMD functions it was discovered that the increase in speed is due to the loading of vector registers. Since in the normal calculation the numbers can be taken directly from memory, this makes it more efficient than SIMD, as SIMD needs to load the numbers from memory into new vector registers to perform the calculation. Profiling indicated that this memory load step took 49.58% of the CPU time for the entire calculation. The actual calculation step only takes 7.5% of the program time. It's clear that the time from this SIMD implementation bottlenecks at the memory load step.

It is clear that the principle of SIMD does same some CPU time for the execution of the poisson algoirthm, but the time taken to set up the calculation outweighs the time saved compared to the normal approach. Thus, the SIMD implementation was not included in the final poisson solver code.

References

[1] I. Advanced Micro Devices, "AMD64 Architecture Programmer's Manual Volume 2: System Programming." 2024. [Online]. Available: https://www.amd.com/system/files/2020-10/amd64-architecture-programmers-manual-volume-2-system-programming.pdf

[2] N. Corporation, "CUDA C Programming Guide." 2023. [Online]. Available: https://docs.nvidia.com/cuda/cuda-c-programming-guide/